

# Architectural Paradigms for Zero-Knowledge, High-Throughput Remote Transfer of Massive Datasets

## The Challenge of Massive Data Transfer in Modern Networks

The contemporary digital landscape is characterized by the exponential growth of data generated by specialized industries, including research laboratories, government agencies, universities, media and visual effects (VFX) studios, software development teams, legal firms, data engineering teams, and healthcare organizations. These entities frequently require the transmission of extremely large, highly sensitive datasets—ranging from 100 gigabytes to several terabytes—across disparate geographical locations. The core problem statement centers on the fact that current file-sharing solutions are fundamentally not optimized for transferring these extremely large, sensitive datasets across long distances while simultaneously maintaining privacy, speed, reliability, and ease of use. The pain point resides not merely in the act of "file sharing," but in overcoming the profound infrastructural bottlenecks associated with massive scale.

Existing infrastructure is fundamentally ill-equipped to handle this specific confluence of requirements, leading to severe operational friction. Traditional cloud storage platforms introduce upload and download bottlenecks, prohibitive storage costs, and rigid file-size restrictions.<sup>1</sup> Standard email and messaging protocols impose strict size limitations, typically capping out at 25 megabytes, which renders them entirely obsolete for modern professional workflows.<sup>3</sup> Conversely, traditional transfer methods and enterprise-grade Managed File Transfer (MFT) systems often demand complex server setups and technical expertise to deploy and maintain.<sup>4</sup> Furthermore, standard transfer protocols force a bifurcation between speed and security; sensitive data must usually be encrypted separately prior to transfer, thereby increasing operational complexity, human error, and the likelihood of security breaches.<sup>5</sup> Adding to the systemic inefficiency, large datasets are often duplicated multiple times across servers, wasting vast amounts of bandwidth and storage space.<sup>6</sup> Secure sharing of large text, documents, datasets, videos, and archives is highly fragmented across different tools, forcing organizations to stitch together disparate solutions to achieve basic operational functionality. The proposed architectural paradigm envisions an "AirDrop for the Internet." This requires building a secure, zero-knowledge platform that enables the encrypted transfer and sharing of massive datasets and files (100GB–TB scale) across long distances through simple shareable links, without the limitations of existing cloud storage solutions. The platform must allow users to send terabytes of encrypted data to anyone, anywhere, using a single secure link—without exposing the data to the platform, sacrificing speed, or managing complex infrastructure. This research report comprehensively analyzes the cryptographic limitations, protocol innovations,

deduplication algorithms, and regulatory compliance frameworks necessary to construct this optimal, high-throughput data transfer architecture.

## **Target Demographics and Industry-Specific Pain Points**

A platform capable of executing zero-knowledge, hardware-agnostic file sharing will rapidly disrupt industries where existing paradigms fail due to regulatory, security, or scale constraints. Analyzing the specific workflows of the target users reveals the critical necessity for a unified, high-throughput solution.

The visual effects (VFX) and media production industry operates on compressed timelines where producers, video editors, colorists, and VFX vendors must collaborate on massive, uncompressed digital assets.<sup>8</sup> During the round-trip workflow, editorial teams require access to the highest fidelity camera originals, as proxy files are insufficient for color grading or CGI compositing.<sup>9</sup> The primary bottleneck is the sheer size of the datasets. A typical untrimmed dailies array can easily exceed 150 terabytes.<sup>11</sup> Due to the exorbitant cost of cloud data transfer and the technical complexities of trimming media for transmission via legacy enterprise tools, studios frequently revert to physical logistics, shipping encrypted hard drives via courier services.<sup>8</sup> A frictionless, link-based platform natively fulfills the strict security mandates of the Trusted Partner Network (TPN) and the Motion Picture Association (MPA), ensuring pre-release content remains immune to interception.<sup>9</sup>

Healthcare organizations require the secure transfer of highly sensitive Protected Health Information (PHI). Medical imaging studies are stored as DICOM files, which embed detailed patient identifiers, diagnostic data, and clinical metadata directly within their headers.<sup>12</sup> A standard radiological study averages 1.5 gigabytes, complete cardiac studies reach 2 gigabytes, and genomic sequencing data routinely exceeds 10 gigabytes.<sup>3</sup> Traditional IT infrastructure in hospitals relies on outdated protocols; staff often use physical CDs, insecure SFTP servers, or risk HIPAA violations by utilizing consumer-grade file-sharing platforms.<sup>3</sup> The financial ramifications of failing to secure this data are catastrophic, with the average healthcare data breach costing \$11.45 million in 2025.<sup>12</sup> A zero-knowledge architecture ensures compliance with HIPAA and the GDPR by encrypting data before it leaves the sender's device, neutralizing the risk of intermediary interception.<sup>3</sup>

Software development and data engineering teams face unique challenges regarding dataset duplication and continuous integration workflows. These teams frequently transfer massive virtual machine (VM) images, compiled binaries, and vast data lakes.<sup>7</sup> When multiple developers upload slightly modified versions of a 50-gigabyte VM image, standard transfer protocols transmit the entire file repeatedly, saturating network bandwidth and exponentially increasing storage costs.<sup>7</sup> A platform optimized for smart data optimization—incorporating deduplication, delta transfers, and chunk optimization—is essential to send only the changed portions of these massive files, accelerating deployment pipelines.<sup>15</sup>

Research labs, universities, government agencies, and legal firms are bound by strict data sovereignty, chain-of-custody, and intellectual property constraints. Legal firms conducting

electronic discovery (e-discovery) must transfer terabytes of indexed documents without risking exposure to opposing counsel or third-party platform operators. Government agencies handling classified or export-controlled data (e.g., ITAR compliance) cannot utilize platforms where the service provider holds the decryption keys.<sup>5</sup> A zero-knowledge architecture, where the server stores encrypted chunks only and the platform never sees plaintext, is the only acceptable paradigm for these highly regulated entities.

## **Economic and Technical Limitations of Existing Solutions**

To understand the necessity for a novel architecture, one must rigorously analyze the technical and economic shortcomings of the current market leaders in large file transfer. The market is currently bisected into self-hosted Managed File Transfer (MFT) solutions and cloud-native, pay-as-you-go platforms, both of which present significant compromises.

Self-hosted and hybrid solutions, such as IBM Aspera and Signiant, are engineered to bypass standard Transmission Control Protocol (TCP) bottlenecks. IBM Aspera utilizes the proprietary Fast Adaptive and Secure Protocol (FASP), a connectionless protocol built upon the User Datagram Protocol (UDP).<sup>17</sup> FASP features a control-theoretic retransmission algorithm that maximizes goodput by utilizing all available network bandwidth, largely unaffected by transfer distance, file size, or network conditions.<sup>17</sup> While technologically robust, these systems present significant Total Cost of Ownership (TCO) challenges. Implementations require specialized software, dedicated servers, manual security patching, and constant IT support.<sup>4</sup> Organizations without the resources to develop and scale complex self-hosted environments find these systems cumbersome, complicated to set up, and time-consuming for onboarding new users.<sup>4</sup> IBM Aspera environments can cost approximately \$12,000 annually for licensing, plus an estimated \$3,500 for server maintenance, resulting in substantial "wasted spend" (estimated at nearly \$1,290 per month) during periods of low transfer volume.<sup>19</sup> Signiant commands similar flat-fee annual subscriptions ranging from \$8,000 to \$15,000.<sup>20</sup>

Conversely, cloud-native solutions like MASV operate on a browser-based delivery model with REST APIs and cloud connectors.<sup>5</sup> MASV eliminates the need for complex self-hosted environments and is capable of accelerating hundreds of gigabytes over the web without requiring file compression or splitting.<sup>4</sup> However, the economic model becomes prohibitive at the enterprise scale. MASV charges per gigabyte; while suitable for small post-production houses transferring under 2 terabytes a month, scaling up to 100 terabytes of data transfer per month costs upwards of \$25,000, scaling linearly thereafter.<sup>5</sup> Furthermore, MASV lacks advanced security controls. Files expire after 7 days, there is no detailed tracking of which buyers reviewed agreements, and there is no protection against unauthorized redistribution, which is a critical failure point for media companies negotiating licensing agreements.<sup>20</sup> Standard consumer cloud platforms (e.g., Google Drive, Dropbox, Proton Drive) suffer from upload and download bottlenecks when handling terabyte-scale data.<sup>2</sup> While some platforms like Proton Drive offer zero-knowledge encryption, users frequently report severe performance issues, staggering, and slow uploads when attempting to push massive files through standard

browser interfaces, indicating underlying systemic limitations in handling high-throughput encrypted streams.<sup>21</sup>

| <b>Solution Architecture</b> | <b>Pricing Model</b>                  | <b>Estimated Annual Cost (At 10 TB/month)</b> | <b>Protocol/Infrastructure</b>       | <b>Key Limitations</b>                                                                                                    |
|------------------------------|---------------------------------------|-----------------------------------------------|--------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| IBM Aspera                   | Enterprise Licensing + Infrastructure | ~\$15,500 (Base) + Hardware                   | UDP-based FASP, Self-hosted/Hybrid   | High setup complexity, requires dedicated IT, wasted spend during slow periods, requires specific ports. <sup>17</sup>    |
| Signiant                     | Flat-fee Annual Subscription          | \$8,000 - \$15,000                            | UDP Acceleration, Self-hosted/Hybrid | High startup costs, complex server maintenance, cumbersome user onboarding, high TCO. <sup>4</sup>                        |
| MASV                         | Pay-as-you-go (per GB)                | ~\$30,000                                     | Cloud-based, Custom Acceleration     | Exorbitant at high volume, files expire after 7 days, lacks detailed tracking, platform has access to data. <sup>20</sup> |
| Consumer Cloud               | Subscription                          | Variable                                      | TCP/HTTP, Cloud-based                | Severe latency over long distances, synchronous                                                                           |

|  |  |  |  |                                                                  |
|--|--|--|--|------------------------------------------------------------------|
|  |  |  |  | processing, size limits, slow encrypted throughput. <sup>4</sup> |
|--|--|--|--|------------------------------------------------------------------|

The data indicates that organizations are forced to choose between exorbitant recurring costs for cloud convenience, high operational complexity for on-premise speed, or severe bottlenecks with standard consumer clouds. A modern architecture must circumvent this dichotomy by providing end-to-end encryption and large-scale remote transfer capabilities without requiring complex infrastructure.

## Cryptographic Bottlenecks in Browser-Based Architectures

Constructing a zero-knowledge, web-based platform requires executing end-to-end encryption directly on the client's device before the data leaves the sender's machine. The platform's servers must only receive and store encrypted chunks, ensuring that the service provider never possesses the cryptographic keys required to access the plaintext, meaning only the intended recipients can decrypt the payload.<sup>6</sup> However, executing cryptographic operations on massive 100GB+ datasets within a standard web browser presents severe systemic challenges, primarily due to the limitations of the native Web Cryptography API. The Web Cryptography API, specifically the SubtleCrypto interface, was designed to provide hardware-accelerated, near-native performance for cryptographic tasks within the browser, utilizing underlying system capabilities like AES-NI instructions.<sup>23</sup> However, it completely lacks support for incremental streaming operations.<sup>25</sup> Mature cryptographic libraries operating on native systems allow developers to initialize a context, feed data incrementally via an `update()` method, and finalize the output, thereby maintaining a constant memory footprint regardless of the file size.<sup>25</sup>

SubtleCrypto demands that data be passed as a single, monolithic `ArrayBuffer`.<sup>25</sup> When attempting to encrypt or hash a 100-gigabyte file, the browser attempts to load the entire dataset into Random Access Memory (RAM), immediately triggering catastrophic memory pressure and inevitable browser tab crashes.<sup>25</sup> Developers have historically attempted various workarounds to bypass this limitation. The "Blob/Response" method leverages other browser APIs to convert data streams into Blob objects and extract the `ArrayBuffer`.<sup>25</sup> This approach is fundamentally flawed as it forces the browser to buffer the entire stream into memory, holding the massive dataset twice during the conversion process, guaranteeing a memory-related crash for large files.<sup>25</sup> Another workaround involves manually collecting chunks into arrays via a stream reader and concatenating them into a giant monolithic buffer.<sup>25</sup> This yields zero memory benefits, blocks the main execution thread, and is highly inefficient.<sup>25</sup>

To bypass SubtleCrypto limitations, sophisticated architectures often rely on compiling audited C or Rust libraries (such as `libsodium`) into WebAssembly (WASM) bundles to handle streaming encryption and modern Key Derivation Functions (KDFs) like `Argon2id`.<sup>25</sup> While WASM executes

with high performance, a critical security vulnerability exists at the boundary between WASM and the JavaScript runtime environment. When a cryptographic key is forged within WASM's linear memory, it must be transferred to SubtleCrypto across the JavaScript boundary as a Uint8Array.<sup>26</sup> During this transfer, the raw key material exists within the JavaScript heap.<sup>26</sup> JavaScript relies on non-deterministic Garbage Collection (GC), meaning the engine does not immediately clear memory upon command.<sup>26</sup> Even if developers proactively write zeroes over the buffers (.fill(0)), internal engine optimizations within V8 or SpiderMonkey may have already duplicated the data, leaving plaintext keys exposed in physical RAM for extended periods.<sup>26</sup> This exposure window severely undermines zero-knowledge guarantees, leaving the system vulnerable to memory dump extractions, Cross-Site Scripting (XSS) attacks, and microarchitectural timing side-channels such as Spectre.<sup>26</sup> There are ongoing initiatives, such as the WinterCG proposal, to integrate streaming processing directly into the Web Cryptography API specification.<sup>27</sup> This proposal advocates extending SubtleCrypto to accept ReadableStream objects for hashing (digest), signing (sign), and verifying (verify), while introducing EncryptionStream and DecryptionStream classes.<sup>27</sup> However, a critical limitation remains regarding Authenticated Encryption with Associated Data (AEAD) ciphers, such as AES-GCM, which is the industry standard for secure file transfer (utilized by platforms like Wormhole for 128-bit AES-GCM streaming encryption).<sup>22</sup> Cryptographic security protocols dictate that authenticated ciphers must never output plaintext until the authentication tag has been fully verified at the end of the stream.<sup>27</sup> Releasing decrypted chunks sequentially before the final tag is processed would expose the application to inauthentic, potentially malicious plaintext.<sup>27</sup> Consequently, the streaming proposal explicitly restricts EncryptionStream and DecryptionStream to unauthenticated ciphers like AES-CBC, rendering them unsuitable for high-security zero-knowledge file transfers without secondary, complex HMAC layering.<sup>24</sup> Furthermore, algorithms like the X448 EdDSA signature scheme require the entire message to be hashed multiple times, making streaming mathematically impossible.<sup>27</sup> Therefore, building an "AirDrop for the Internet" requires hybrid chunking mechanisms and Service Workers to encrypt streams efficiently without waiting for native streaming APIs to mature.<sup>28</sup>

## Semantic Security vs. Deduplication: Message-Locked Encryption

A platform transferring datasets at the terabyte scale must optimize bandwidth and storage utilization to remain economically viable. Large datasets, especially in software development and media, are frequently duplicated or undergo minor revisions.<sup>7</sup> Smart data optimization requires deduplication, which identifies identical files or file chunks across users and stores only a single instance.<sup>6</sup>

However, a fundamental technical contradiction exists between end-to-end encryption and deduplication. Standard semantic security ensures that identical plaintexts encrypted with unique, randomly generated keys or initialization vectors will produce entirely different ciphertexts.<sup>29</sup> Because the zero-knowledge server is blind to the plaintext and sees only

randomized ciphertext, it cannot identify duplicate files uploaded by different users.<sup>6</sup> If a team of developers uploads the same 50GB file encrypted with different keys, the server stores 50GB multiple times, completely negating the bandwidth and storage benefits of cloud architecture.

To resolve this tension, advanced systems utilize Message-Locked Encryption (MLE), the most prominent manifestation of which is Convergent Encryption.<sup>6</sup> In MLE, the cryptographic encryption key is deterministically derived from the content of the message itself (e.g., computing the SHA-256 hash of the plaintext file and using that hash as the symmetric encryption key).<sup>29</sup> Because identical files will yield identical hashes, they will be encrypted into identical ciphertexts, allowing the server to recognize the duplicate blobs and perform deduplication over encrypted data without ever knowing the underlying plaintext.<sup>29</sup>

The critical vulnerability of standard Convergent Encryption is that it is inherently subject to offline brute-force dictionary attacks.<sup>6</sup> If an attacker has a known set of possible files, they can independently compute the hash, generate the encryption key, encrypt the guessed file, and check if the resulting ciphertext exists on the server, thereby revealing whether a target user possesses that specific file.<sup>6</sup> To achieve strong confidentiality while preserving deduplication, architectures such as DupLESS enforce server-aided encryption.<sup>6</sup> In this architecture, clients do not generate message-based keys in isolation. Instead, they interact with a dedicated key-server via an oblivious Pseudorandom Function (PRF) protocol to obtain the keys.<sup>6</sup> This PRF protocol ensures that the key-server assists in generating the key without learning the file's hash, and the client cannot perform brute-force checks without querying the key-server, which can enforce rate limiting. This enables clients to store encrypted data, have the service perform deduplication on their behalf, and maintain robust confidentiality guarantees that resist brute-force attacks.<sup>6</sup>

## **Data Optimization: Content-Defined Chunking and Delta Transfers**

Deduplicating massive files is highly inefficient if performed exclusively at the macro-file level, as a single byte modification alters the entire file hash, requiring the re-transmission of the entire terabyte dataset. To achieve true smart data optimization, the platform must utilize a delta-transfer algorithm, segmenting files into chunks and transmitting only the modified blocks.<sup>15</sup> Standard synchronization utilities, such as rsync, utilize a sliding window to match data chunks byte-by-byte.<sup>7</sup> However, rsync employs invariant chunking and compression methods, making it computationally heavy and unable to adapt to varying network environments.<sup>16</sup> Furthermore, traditional fixed-length chunking suffers from the byte-shifting problem; inserting a single byte at the beginning of a file shifts all subsequent boundaries, altering the hashes of every subsequent chunk and negating redundancy detection.<sup>33</sup>

Modern, bandwidth-efficient synchronization requires Content-Defined Chunking (CDC).<sup>7</sup> CDC determines chunk boundaries dynamically based on the internal data features rather than fixed byte offsets.<sup>33</sup> This is mathematically achieved using a rolling hash, most commonly the Rabin fingerprint.<sup>34</sup> The Rabin fingerprint is a polynomial rolling hash that treats sequences of bytes as

coefficients over a Galois field.<sup>34</sup> It offers  $O(1)$  sliding updates; as the algorithm slides a search window byte-by-byte across the file, the hash is updated continuously based on minimum, average, and maximum chunk size parameters (e.g., 4 KB, 8 KB, and 16 KB).<sup>7</sup> When the rolling hash matches a specific pre-determined bitmask, a chunk boundary is declared.<sup>35</sup> This deterministic behavior ensures that if data is inserted, edited, or deleted, only the immediate chunk containing the mutation is altered, while all subsequent chunk boundaries rapidly resynchronize, preserving the deduplication ratios of the rest of the file.<sup>34</sup> To support lightweight delta synchronization for high-bandwidth cloud storage, academic and industry researchers have proposed several optimizations to CDC. The FastCDC algorithm replaces the CPU-intensive Rabin fingerprinting with precomputed Gear tables generated at compile time, drastically reducing the nanoseconds per operation and increasing throughput.<sup>36</sup> Frameworks like Dsync and NetSync utilize a fast weak hash called FastFP piggybacked on CDC rolling hashes.<sup>7</sup> NetSync further minimizes sync time by adaptively choosing chunking parameters and compression methods according to current network conditions, performing 2x to 10x faster than traditional rsync models.<sup>16</sup> Integrating these chunking algorithms with Message-Locked Encryption creates highly efficient hybrid schemes like FASTSync and FeatureSync.<sup>37</sup> These feature-based delta sync schemes are tailored specifically for high-bandwidth, encrypted cloud storage environments, reducing cloud sync time by up to 72.6% and cloud sync traffic by 78.5% compared to state-of-the-art sync schemes.<sup>40</sup> By running WASM-accelerated FastCDC algorithms combined with MLE on the client side, the platform ensures parallel chunk-based transmission where only the deduplicated, unique, and compressed encrypted chunks of massive datasets are transmitted over the network, drastically reducing bandwidth consumption and enabling the capability to resume interrupted transfers seamlessly.

## Transport Layer Revolution: Overcoming TCP with WebTransport and QUIC

Regardless of cryptographic efficiency and smart chunk optimization, the physical transmission of terabytes of data across global distances remains constrained by the underlying transport protocols. The Transmission Control Protocol (TCP), which underpins legacy WebSockets and standard HTTP traffic, guarantees reliable, ordered delivery.<sup>41</sup> However, TCP's congestion control mechanisms react to packet loss by drastically halving the transmission window.<sup>42</sup> On high-latency, long-distance Wide Area Networks (WANs), natural packet loss forces TCP to perpetually throttle bandwidth, resulting in severe performance degradation and an inability to saturate available network pipes.<sup>18</sup>

To bypass TCP's shortcomings, enterprise MFT systems like IBM Aspera rely on connectionless UDP frameworks, specifically the FASP protocol.<sup>18</sup> UDP does not inherently mandate handshakes or wait for acknowledgment receipts, avoiding the overhead of TCP.<sup>18</sup> FASP improves upon raw UDP by implementing adaptive rate control and precise, control-theoretic retransmission algorithms, ensuring that only genuinely lost packets are requested by the

recipient.<sup>18</sup> This allows FASP to maintain maximum throughput irrespective of geographical distance or high packet loss rates.<sup>18</sup> The fatal flaw of this approach for a frictionless platform is that it requires specialized client software, dedicated server ports, firewall configurations, and cannot run natively within a web browser without plugins.<sup>4</sup>

For a browser-native platform to achieve MFT-level speeds without requiring installed software, the architecture must transition from WebSockets to WebTransport. WebTransport is a versatile, low-level browser API that operates over HTTP/3, utilizing the QUIC protocol.<sup>41</sup> QUIC is a modern transport protocol built on top of UDP that integrates connection establishment and TLS 1.3 cryptographic handshakes into a single, highly efficient process.<sup>42</sup> WebTransport provides several critical architectural advantages over legacy WebSockets for massive file transfers:

1. **Elimination of Head-of-Line Blocking:** In TCP, if a single packet is lost, the entire stream of subsequent packets is halted until the lost packet is retransmitted.<sup>48</sup> QUIC establishes multiple independent logical streams within a single connection.<sup>42</sup> If a packet belonging to one stream is lost, only that specific stream is paused; the other streams continue to transmit data unimpeded.<sup>42</sup> This multiplexing is vital for parallel, chunk-based transmission of 100GB+ files, allowing efficient file transfer and reducing the overhead of establishing new connections.<sup>50</sup>
2. **Unreliable Datagrams and Reliable Streams:** WebTransport allows developers to route data via bidirectional reliable streams or unreliable datagrams.<sup>46</sup> The platform can send control metadata (e.g., chunk hashes, structural manifests) via reliable streams while blasting encrypted file chunks via unsequenced datagrams, reassembling and verifying them at the application layer, mimicking the performance profile of Aspera's FASP natively in the browser.<sup>42</sup>
3. **Network Resilience and Resumability:** Unlike TCP, which drops connections during network transitions (e.g., switching from Wi-Fi to a cellular network), QUIC utilizes connection IDs that survive IP address changes.<sup>41</sup> This allows massive transfers to resume seamlessly without initiating full reconnects, fulfilling the requirement to easily resume interrupted transfers.
4. **Hardware Independence:** WebTransport is hardware independent, supported by modern browsers, and navigates NATs and firewalls natively by operating over standard HTTPS/UDP ports, eliminating the complex infrastructure requirements of UDP-based MFT tools.<sup>46</sup>

Empirical analyses indicate that replacing TCP/WebSockets with QUIC-based WebTransport significantly reduces latency, optimizes resource utilization, and allows servers to handle a larger number of simultaneous connections without compromising performance.<sup>49</sup>

## Secure Share Links and Access Control Mechanics

A defining characteristic of the proposed solution is moving away from traditional fragmented file sharing mechanisms to the generation of Secure Share Links. Instead of attaching files directly or granting access to a persistent cloud drive folder, the platform generates encrypted links. This UX paradigm transforms the complex underlying cryptography and transport

protocols into a simple, user-friendly interface.

In a zero-knowledge architecture, the decryption key must never be transmitted to the server. To achieve this via shareable links, the platform generates the encryption key locally on the sender's device. After the file is encrypted and chunked, the ciphertext is uploaded to the server via WebTransport. The decryption key is then appended to the generated URL as a fragment identifier (the portion of the URL following the # symbol). Because modern web browsers do not transmit the URL fragment to the server during an HTTP request, the platform's backend remains entirely blind to the key. When the recipient clicks the link, the client-side application extracts the key from the fragment, downloads the encrypted chunks, and decrypts the file locally in the browser.

To further enforce security for sensitive enterprise and government data, these Secure Share Links must support granular access controls. The platform must allow senders to set optional passwords, which act as a secondary Key Encryption Key (KEK) to wrap the primary decryption key, ensuring that even if the link is intercepted, the data remains inaccessible without the out-of-band password. Furthermore, administrators can enforce expiry times and one-time access policies. If a link is set to one-time access, the server deletes the ciphertext chunks immediately upon the first successful download confirmation, mathematically guaranteeing that the data cannot be subsequently accessed or redistributed, a critical feature for legal discovery and media licensing workflows.<sup>20</sup>

## Global Regulatory Compliance and Data Sovereignty

Beyond technical constraints, a global file transfer platform must navigate an increasingly hostile regulatory landscape characterized by data sovereignty, localization mandates, and aggressive incident reporting frameworks. Cross-border data transfers are heavily scrutinized, creating compliance minefields for multinational organizations.<sup>52</sup>

India's Digital Personal Data Protection Act (DPDPA), enacted in 2023, represents a significant evolution in global privacy law. Under Section 16 of the DPDPA, the Indian Central Government retains the explicit power to restrict the transfer of personal data to notified jurisdictions (the "blacklist" approach), diverging from the GDPR's "whitelist" adequacy decision model.<sup>52</sup>

Furthermore, sector-specific regulators enforce stringent localization rules requiring payment, trading, and insurance data to remain exclusively within Indian territory.<sup>56</sup> The DPDPA assigns heavy liabilities, with penalties reaching up to 250 Crore INR for Data Fiduciaries that fail to implement robust technical safeguards.<sup>58</sup>

For organizations routing multi-gigabyte datasets globally, these regulations necessitate intricate legal arrangements. However, the deployment of end-to-end encryption and zero-knowledge architectures functions as a powerful compliance mechanism.<sup>59</sup> Because zero-knowledge platforms encrypt the data client-side and the server operator is technologically incapable of decrypting the payload, the platform acts as a blind conduit.<sup>57</sup> This architectural design fundamentally shifts the data residency model. The regulatory conversation transitions from "where is the data geographically stored?" to "who holds the cryptographic keys to access the data?".<sup>53</sup> Consequently, zero-knowledge infrastructure significantly reduces the compliance burden on Data Fiduciaries, enabling global collaboration

without violating territorial data access strictures and serving as a technical safe harbor. In addition to privacy laws, operational infrastructure is heavily governed by cybersecurity directives. The Indian Computer Emergency Response Team (CERT-In) issued directions under Section 70B of the Information Technology Act, imposing aggressive monitoring and reporting obligations on service providers, cloud operators, and Virtual Private Network (VPN) providers.<sup>61</sup>

Key mandates include:

1. **6-Hour Incident Reporting:** Organizations must report qualified cybersecurity incidents to CERT-In within six hours of discovery or being brought to notice.<sup>61</sup> This aggressive timeframe requires treating the six-hour window as a rapid coordination phase rather than a comprehensive root-cause investigation period.<sup>64</sup>
2. **Log Retention:** Entities must synchronize their system clocks to the National Informatics Centre (NIC) Network Time Protocol (NTP) servers and maintain exhaustive operational logs (firewall, event, application, and VPN logs) for a rolling period of 180 days within India.<sup>61</sup>
3. **Extended KYC Requirements:** Cloud service providers, virtual private servers (VPS), and VPN services are required to collect and store extensive Know-Your-Customer (KYC) details—including validated names, IPs, email addresses, and ownership patterns—for a minimum of five years after account closure.<sup>65</sup>

These mandates effectively ban "anonymous" infrastructural routing within the jurisdiction.<sup>67</sup> For a platform providing encrypted data routing, absolute compliance requires rigorous internal metadata logging decoupled from the encrypted payloads. The platform must maintain logs of connection IPs, transfer volumes, and timestamps for 180 days to satisfy CERT-In, while simultaneously utilizing zero-knowledge encryption to guarantee that the actual file contents remain mathematically inaccessible to both the platform operators and any intercepting authorities.

## Synthesis and Strategic Recommendations

To effectively build the conceptual "AirDrop for the Internet"—a platform capable of routing 100GB to terabyte-scale sensitive datasets globally via simple shareable links—the architecture must meticulously integrate multiple bleeding-edge technologies to overcome the physical, cryptographic, and regulatory failures of existing platforms.

1. **Client-Side Processing Execution:** Due to the severe memory limitations of the native Web Cryptography API (SubtleCrypto) in handling streaming monolithic buffers, and the security vulnerabilities of the JavaScript heap, the client architecture must utilize Service Workers and WebAssembly (WASM). WASM will execute Content-Defined Chunking (CDC) via FastCDC to segment massive files into variable-sized chunks resistant to byte-shifting, without crashing the browser's memory footprint.<sup>28</sup>
2. **Zero-Knowledge Deduplication:** Once chunked, the client must apply Message-Locked Encryption (MLE) via an oblivious PRF key-server model like DupLESS.<sup>6</sup> This protocol derives the cryptographic key from the chunk's content securely, allowing the central zero-knowledge server to safely deduplicate identical encrypted chunks uploaded by

different users (enabling smart delta transfers) without compromising semantic security or exposing the system to brute-force dictionary attacks.

3. **Transport Protocol Revolution:** The transmission of these encrypted, deduplicated, and compressed chunks must abandon legacy TCP and WebSockets. The system must leverage WebTransport over HTTP/3 (QUIC).<sup>41</sup> By utilizing QUIC's multiplexed streams and datagrams, the platform can blast encrypted chunks in parallel, eliminating head-of-line blocking and saturating the network pipe regardless of the geographic distance, mimicking enterprise FASP speeds natively in the browser.<sup>18</sup>
4. **UX and Access Control:** The platform must abstract these complex operations behind Secure Share Links, where decryption keys are passed securely via URL fragments. Senders must have the capability to append optional passwords, set expiry times, and enforce one-time access destruction policies to maintain strict control over their distributed assets.
5. **Regulatory Resilience:** By maintaining strict zero-knowledge architecture, the platform insulates itself and its enterprise clients (in healthcare, VFX, legal, and government) from the catastrophic liabilities of data breaches and HIPAA/TPN violations.<sup>3</sup> Furthermore, it inherently satisfies strict cross-border data transfer limitations enforced by frameworks like the DPDPA, while parallel logging infrastructure ensures compliance with aggressive cybersecurity mandates like CERT-In's 6-hour reporting window.<sup>52</sup>

By synthesizing WASM-driven CDC, Message-Locked Encryption, and QUIC-based WebTransport into a unified, link-based workflow, the resulting platform will deliver unprecedented speed, security, and storage optimization. This architecture fundamentally obsoletes expensive legacy MFT servers, cumbersome VPN setups, and inadequate consumer cloud storage models, providing a definitive solution for massive-scale, zero-knowledge data transfer.

## Works cited

1. We rebuilt Proton Drive's engine. Now it's up to 3x faster on all platforms. : r/ProtonDrive, accessed on June 7, 2026, [https://www.reddit.com/r/ProtonDrive/comments/1twkbdj/we\\_rebuilt\\_proton\\_drives\\_engine\\_now\\_its\\_up\\_to\\_3x/](https://www.reddit.com/r/ProtonDrive/comments/1twkbdj/we_rebuilt_proton_drives_engine_now_its_up_to_3x/)
2. ProtonDrive maximum file size? - Reddit, accessed on June 7, 2026, [https://www.reddit.com/r/ProtonDrive/comments/1gis7fa/protondrive\\_maximum\\_file\\_size/](https://www.reddit.com/r/ProtonDrive/comments/1gis7fa/protondrive_maximum_file_size/)
3. How to Share Large Files Efficiently, Securely, and in Compliance With HIPAA: A Guide for Healthcare Providers - Kiteworks, accessed on June 7, 2026, <https://www.kiteworks.com/secure-file-sharing/large-file-sharing-healthcare/>
4. Signiant Vs Aspera: There Is A Better Solution - MASV, accessed on June 7, 2026, <https://massive.io/alternatives/signiant-vs-aspera/>
5. Aspera vs MASV: The Complete Comparison Guide for Enterprise File Transfer, accessed on June 7, 2026, <https://pacgenesis.com/aspera-vs-masv-enterprise-file-transfer/>

6. DupLESS: Server-Aided Encryption for Deduplicated Storage ..., accessed on June 7, 2026,  
<https://www.usenix.org/conference/usenixsecurity13/technical-sessions/presentation/bellare>
7. Dsync: a Lightweight Delta Synchronization Approach for Cloud Storage Services - Lingfeng Xiang, accessed on June 7, 2026,  
<https://lingfenghsiang.github.io/docs/DSync.pdf>
8. VFX Pipeline: A Complete Guide For Video & Media Pros - MASV, accessed on June 7, 2026, <https://massive.io/workflow/vfx-pipeline/>
9. Video Post-Production Workflow Guide | Frame.io, accessed on June 7, 2026,  
<https://workflow.frame.io/guide/>
10. 4 of Today's Biggest Media Management Challenges. Solved. - Signiant, accessed on June 7, 2026,  
<https://www.signiant.com/resources/articles/4-media-management-challenges/>
11. Virtual Roundtable: Storage - postPerspective, accessed on June 7, 2026,  
<https://postperspective.com/virtual-roundtable-storage-4/>
12. Secure DICOM Sharing: Best Practices for Protecting Patient Data - OmniPACS, accessed on June 7, 2026,  
<https://www.omnipacs.com/secure-dicom-sharing-best-practices-for-protecting-patient-data/>
13. HIPAA Compliant File Sharing: 3 Major Risks, accessed on June 7, 2026,  
<https://www.hipaavault.com/sftp/3-file-sharing-risks-in-healthcare/>
14. Image Data Sharing for Biomedical Research—Meeting HIPAA Requirements for De-identification - PMC, accessed on June 7, 2026,  
<https://pmc.ncbi.nlm.nih.gov/articles/PMC3264712/>
15. rsync is the only backup tool you actually need - XDA Developers, accessed on June 7, 2026,  
<https://www.xda-developers.com/rsync-only-backup-tool-you-actually-need/>
16. NetSync: A Network Adaptive and Deduplication-Inspired Delta Synchronization Approach for Cloud Storage Services - Group, accessed on June 7, 2026,  
<https://www.thucloud.com/zhenhua/papers/TPDS'22%20NetSync.pdf>
17. IBM Aspera vs. MASV Comparison - SourceForge, accessed on June 7, 2026,  
<https://sourceforge.net/software/compare/IBM-Aspera-vs-MASV/>
18. Fast and Secure Protocol - Wikipedia, accessed on June 7, 2026,  
[https://en.wikipedia.org/wiki/Fast\\_and\\_Secure\\_Protocol](https://en.wikipedia.org/wiki/Fast_and_Secure_Protocol)
19. Aspera Alternative At 1/4 The Cost - MASV, accessed on June 7, 2026,  
<https://massive.io/alternatives/aspera-alternative/>
20. My Honest Review of MASV Alternatives (\$0.25/GB Adds Up Fast) in 2026 — Peony, accessed on June 7, 2026, <https://www.peony.ink/blog/masv-alternatives>
21. ProtonDrive Super slow upload - Constant staggering : r/ProtonMail - Reddit, accessed on June 7, 2026,  
[https://www.reddit.com/r/ProtonMail/comments/wr2l5y/protondrive\\_super\\_slow\\_upload\\_constant\\_staggering/](https://www.reddit.com/r/ProtonMail/comments/wr2l5y/protondrive_super_slow_upload_constant_staggering/)
22. Security Design - Wormhole, accessed on June 7, 2026,  
<https://wormhole.app/security>

23. Web Crypto API - MDN Web Docs, accessed on June 7, 2026,  
[https://developer.mozilla.org/en-US/docs/Web/API/Web\\_Crypto\\_API](https://developer.mozilla.org/en-US/docs/Web/API/Web_Crypto_API)
24. Comparing Performance of JavaScript Cryptography Libraries | by Encryb - Medium, accessed on June 7, 2026,  
<https://medium.com/@encryb/comparing-performance-of-javascript-cryptography-libraries-42fb138116f3>
25. Web Crypto's SubtleCrypto: A Masterclass in Developer Hostility and ..., accessed on June 7, 2026,  
<https://misakikasumi.medium.com/web-cryptos-subtlecrypto-a-masterclass-in-developer-hostility-and-how-it-strangles-the-modern-web-76d1748ceef>
26. The Memory Gap in WASM-to-WebCrypto Bridges. And how to build ..., accessed on June 7, 2026,  
[https://www.reddit.com/r/rust/comments/1q44ey2/the\\_memory\\_gap\\_in\\_wasmtowebcrypto\\_bridges\\_and\\_how/](https://www.reddit.com/r/rust/comments/1q44ey2/the_memory_gap_in_wasmtowebcrypto_bridges_and_how/)
27. proposal-webcrypto-streams/explainer.md at main · WinterTC55 ..., accessed on June 7, 2026,  
<https://github.com/wintercg/proposal-webcrypto-streams/blob/main/explainer.md>
28. memory efficient in-browser large file encryption using libsodium ..., accessed on June 7, 2026,  
[https://www.reddit.com/r/crypto/comments/ntjft9/memory\\_efficient\\_inbrowser\\_large\\_file\\_encryption/](https://www.reddit.com/r/crypto/comments/ntjft9/memory_efficient_inbrowser_large_file_encryption/)
29. Is there a form of cryptography where the key is derived from the plaintext, accessed on June 7, 2026,  
<https://crypto.stackexchange.com/questions/93635/is-there-a-form-of-cryptography-where-the-key-is-derived-from-the-plaintext>
30. Secure Deduplication with Efficient and Reliable Convergent Key Management - CNSR@VT, accessed on June 7, 2026,  
<https://www.cnsr.ictas.vt.edu/publication/t1.pdf>
31. SimEnc: A High-Performance Similarity-Preserving Encryption Approach for Deduplication of Encrypted Docker Images - USENIX, accessed on June 7, 2026,  
<https://www.usenix.org/system/files/atc24-sun.pdf>
32. rsync Over SSH: Secure, Fast File Transfers and Backups on Linux | LinuxBlog.io, accessed on June 7, 2026, <https://linuxblog.io/rsync-over-ssh/>
33. Delta Encoding Overhead Analysis of Cloud Storage Systems Using Client-Side Encryption, accessed on June 7, 2026,  
[https://www.researchgate.net/publication/338941770\\_Delta\\_Encoding\\_Overhead\\_Analysis\\_of\\_Cloud\\_Storage\\_Systems\\_Using\\_Client-Side\\_Encryption](https://www.researchgate.net/publication/338941770_Delta_Encoding_Overhead_Analysis_of_Cloud_Storage_Systems_Using_Client-Side_Encryption)
34. System Design: Dropbox (File Sync, Chunking, and Deduplication ..., accessed on June 7, 2026, <https://crackingwalnuts.com/post/dropbox-system-design>
35. Coupling Secret Sharing with Decentralized Server-Aided Encryption in Encrypted Deduplication - MDPI, accessed on June 7, 2026,  
<https://www.mdpi.com/2076-3417/15/3/1245>
36. Introducing go-cdc-chunkers: chunk and deduplicate everything - Plakar, accessed on June 7, 2026,

- <https://plakar.io/posts/2025-07-11/introducing-go-cdc-chunkers-chunk-and-deduplicate-everything/>
37. Selected Publications - 厦门大学先进存储技术实验室, accessed on June 7, 2026, <https://astl.xmu.edu.cn/Publications.htm>
  38. Zhirong Shen, accessed on June 7, 2026, <https://shenzr.github.io/>
  39. SkySync: Accelerating File Synchronization with Collaborative Delta Generation - USENIX, accessed on June 7, 2026, [https://www.usenix.org/system/files/conf%C3%A9rence/fast26/fast26spring-prep\\_ub\\_zhang-zhihao.pdf](https://www.usenix.org/system/files/conf%C3%A9rence/fast26/fast26spring-prep_ub_zhang-zhihao.pdf)
  40. When Delta Sync Meets Message-Locked Encryption: a Feature-based Delta Sync Scheme for Encrypted Cloud Storage - IEEE Xplore, accessed on June 7, 2026, <https://ieeexplore.ieee.org/document/9546438/>
  41. Building Real-Time Web Apps with WebTransport (Replacing WebSockets?), accessed on June 7, 2026, <https://dev.to/mukhilpadmanabhan/building-real-time-web-apps-with-webtransport-replacing-websockets-3348>
  42. (PDF) Analysis of Next-Generation Internet Transport Protocols: QUIC, WebTransport, HTTP/3 - ResearchGate, accessed on June 7, 2026, [https://www.researchgate.net/publication/399126043\\_Analysis\\_of\\_Next-Generation\\_Internet\\_Transport\\_Protocols\\_QUIC\\_WebTransport\\_HTTP3](https://www.researchgate.net/publication/399126043_Analysis_of_Next-Generation_Internet_Transport_Protocols_QUIC_WebTransport_HTTP3)
  43. IBM Aspera and FASP Protocol - Dot Group, accessed on June 7, 2026, <https://dotgroup.ai/2024/02/27/ibm-aspera-and-fasp-protocol/>
  44. What is Aspera FASP®? Explaining the FASP Protocol. - PacGenesis, accessed on June 7, 2026, <https://pacgenesis.com/what-is-aspera-fasp/>
  45. Using FASP - IBM, accessed on June 7, 2026, <https://www.ibm.com/docs/en/connect-direct/6.2.0?topic=v62-using-fasp>
  46. WebSockets vs Server-Sent-Events vs Long-Polling vs WebRTC vs WebTransport - RxDB, accessed on June 7, 2026, <https://rxdb.info/articles/websockets-sse-polling-webrtc-webtransport.html>
  47. WebTransport is almost here to allow UDP-like exchange in the browser | Hacker News, accessed on June 7, 2026, <https://news.ycombinator.com/item?id=45820782>
  48. What is WebTransport and can it replace WebSockets? - Ably Realtime, accessed on June 7, 2026, <https://ably.com/blog/can-webtransport-replace-websockets>
  49. WebTransport: The Future of Real-Time Communication, Bridging the Gap Beyond WebRTC & Websockets | by Lakin Mohapatra, accessed on June 7, 2026, <https://lakin-mohapatra.medium.com/web-transport-the-future-of-real-time-communication-bridging-the-gap-beyond-webrtc-e1203f284ae9>
  50. What is WebTransport and How Does it Work? - PubNub, accessed on June 7, 2026, <https://www.pubnub.com/guides/webtransport/>
  51. network abstractions and emulation for distributed systems - ASC - UNL, accessed on June 7, 2026, <https://asc.di.fct.unl.pt/~jleitao/thesis/DanielJoaoMSc.pdf>
  52. Guidance on Cross-Border Data Transfers for Indian Organisations - OneTrust DataGuidance, accessed on June 7, 2026,

- [https://www.dataguidance.com/sites/default/files/dcsi\\_privacy\\_across\\_borders\\_guidance\\_on\\_cross-border\\_data\\_transfers\\_for\\_indian\\_organizations.pdf](https://www.dataguidance.com/sites/default/files/dcsi_privacy_across_borders_guidance_on_cross-border_data_transfers_for_indian_organizations.pdf)
53. The Global Rise of Data Localization: Risks, Tradeoffs, and What Comes Next - TrustArc, accessed on June 7, 2026,  
<https://trustarc.com/resource/global-rise-data-localization-risks/>
  54. Digital Personal Data Protection Act, 2023 DPDPA SECTION 16 WITH INTERPRETATION, accessed on June 7, 2026,  
<https://www.dpdpa.com/dpdpa2023/chapter-4/section16.html>
  55. India's Digital Personal Data Protection Act 2023 vs. the GDPR: A Comparison - Latham & Watkins, accessed on June 7, 2026,  
<https://www.lw.com/admin/upload/SiteAttachments/Indias-Digital-Personal-Data-Protection-Act-2023-vs-the-GDPR-A-Comparison.pdf>
  56. Cross-Border Data Transfers After DPDP: A Practical Guide - Digio, accessed on June 7, 2026,  
<https://www.digio.in/blog/cross-border-data-transfers-after-dpdp-a-practical-guide-2/>
  57. Hybrid Sovereignty: Building Split-Brain Databases via Secure Tunnels - Medium, accessed on June 7, 2026,  
<https://medium.com/@instatunnel/hybrid-sovereignty-building-split-brain-databases-via-secure-tunnels-a6d2fef2d490>
  58. Designing Enterprise Privacy Journey For End-to-End DPDP Compliance, accessed on June 7, 2026,  
[https://cdn.prod.website-files.com/650dc6e8c3edc32394341ca0/694290b312976f49c6858dea\\_DPDP%20White%20paper%203\\_compressed.pdf](https://cdn.prod.website-files.com/650dc6e8c3edc32394341ca0/694290b312976f49c6858dea_DPDP%20White%20paper%203_compressed.pdf)
  59. The Global Data Residency Crisis: How Enterprises Can Navigate Geolocation, Storage, and Privacy Compliance Without Sacrificing Performance - Deepak Gupta, accessed on June 7, 2026,  
<https://guptadeepak.com/the-global-data-residency-crisis-how-enterprises-can-navigate-geolocation-storage-and-privacy-compliance-without-sacrificing-performance/>
  60. Tech Lawyer's Webmail Guide 2026 — Gmail, Outlook, Zoho, Proton, DeoMail | Advocate Siddharth Gupta, accessed on June 7, 2026,  
<https://siddharthgupta.in/blog/tech-lawyer-webmail-guide-gmail-outlook-zoho-deomail-2026>
  61. CERT-IN Directions - AZB & Partners, accessed on June 7, 2026,  
<https://www.azbpartners.com/bank/cert-in-directions/>
  62. Everything You Need To Know About India's New Guidelines Related to Cyber Incident Reporting by CERT-In - Ankura.com, accessed on June 7, 2026,  
<https://ankura.com/insights/everything-you-need-to-know-about-indias-new-guidelines-related-to-cyber-incident-reporting-by-cert-in>
  63. India - A new and onerous cyber security framework, with breach reporting obligations, accessed on June 7, 2026,  
<https://www.linklaters.com/en/insights/blogs/digilinks/2022/may/india-new-and-onerous-cyber-security-framework-and-breach-reporting-obligations>
  64. CERT-In Incident Reporting: 6-Hour Rule and Log Retention Best Practices,

accessed on June 7, 2026,

<https://www.atrity.com/cert-in-incident-reporting-6-hour-rule-and-log-retention-best-practices/>

65. GreyNoise | New reporting requirements from Cert-In 70B, accessed on June 7, 2026,

<https://www.greynoise.io/blog/comply-with-cert-ins-new-reporting-requirements-by-cutting-irrelevant-alerts>

66. CERT-In Directive - A Step to Strengthen India's Cybersecurity Posture - SISA, accessed on June 7, 2026,

<https://sisa.ai/resource/blog/cert-in-directive-a-step-to-strengthen-indias-cybersecurity-posture>

67. CERT-In Compliance Guide: The 6-Hour Rule & How to Avoid Penalties - MailArmor AI, accessed on June 7, 2026,

<https://mailarmor.ai/cert-in-compliance-guide>

68. How Do Organizations Comply With the 2022 CERT-In Directions? - InfoSec Brigade - Information and Cyber Security Solution, accessed on June 7, 2026,

<https://infosecbrigade.com/how-do-organizations-comply-with-the-2022-cert-in-directions/>